

Lab 2 Solution - GP Regression in BoTorch

July 1, 2022

1 Lab 2 Solution: Gaussian process regression in BoTorch and gpytorch


BoTorch is a package for Bayesian Optimization. Its support for Gaussian process regression comes from another package, gpytorch. Both of these packages are, in turn, based on pytorch, which is a framework often used for deep learning that supports fast operations on GPUs in python.

Here we'll use these packages to do the same inference that we did above, but more quickly and with more features. The format for this portion of the tutorial is based on the BoTorch tutorial, https://botorch.org/tutorials/fit_model_with_torch_optimizer, which in turn seems to be based on the gpytorch tutorial, https://github.com/cornellius-gp/gpytorch/blob/master/examples/01_Exact_GPs/Simple_GP_Regression.ipynb

```
[1]: import torch # loading torch before matplotlib can be important
import matplotlib.pyplot as plt
import math
import numpy as np
import os.path
```

```
[2]: filename = 'lab1_data.csv'
data = np.loadtxt(filename)
train_x = data[:,0] # First column of the data
train_y = data[:,1] # Second column of the data
plt.xlim(0,1)
plt.plot(train_x, train_y, 'o-')
```

```
[2]: [<matplotlib.lines.Line2D at 0x7ff790652e50>]
```




```
[3]: # use a GPU if available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
dtype = torch.double

[4]: # Read in data from a file. Code copied from Lab 1.
filename = 'lab1_data.csv'

# If data doesn't exist, generate it
if ~os.path.exists(filename):
    np.random.seed(1)
    train_x = np.sort(np.random.rand(10)) # 10 points, uniformly distributed
    ↪ between 0 and 1
    train_y = [math.sin(x * (2 * math.pi)) + 0.0 * np.random.randn() for x in
    ↪ train_x]
    data = np.transpose([train_x, train_y])
    np.savetxt(filename, data)

# Read in data from a file.
data = np.loadtxt(filename)
train_x = data[:,0] # First column of the data
train_y = data[:,1] # Second column of the data
plt.xlim(0,1)
plt.plot(train_x, train_y, 'o-')
```

```
[4]: [<matplotlib.lines.Line2D at 0x7ff760b169d0>]
```



1.0.1 Convert training data to torch format

```
[5]: # Convert our training x and y data to torch format.
# To do this, we first convert to a torch tensor.
# Then we use unsqueeze to make the tensors multidimensional.
# Note that we are changing the datatype of train_x and train_y
train_x = torch.from_numpy(train_x).unsqueeze(1)
train_y = torch.from_numpy(train_y).unsqueeze(1)

# Outputting one of these tensors to show what they look like
train_x
```

```
[5]: tensor([[1.1437e-04],
          [9.2339e-02],
          [1.4676e-01],
          [1.8626e-01],
          [3.0233e-01],
          [3.4556e-01],
          [3.9677e-01],
          [4.1702e-01],
          [5.3882e-01],
          [7.2032e-01]], dtype=torch.float64)
```

1.0.2 Train the GP regression model and plot the posterior

We will model the function using a `SingleTaskGP`, which by default uses a `GaussianLikelihood` and infers the unknown noise level.

The default optimizer for the `SingleTaskGP` is L-BFGS-B, which takes as input explicit bounds on the noise parameter. However, the `torch` optimizers don't support parameter bounds as input. To use the `torch` optimizers, then, we'll need to manually register a constraint on the noise level. When registering a constraint, the `softplus` transform is applied by default, enabling us to enforce a lower bound on the noise.

Note: Without manual registration, the model itself does not apply any constraints, due to the interaction between constraints and transforms. Although the `SingleTaskGP` constructor does in fact define a constraint, the constructor sets `transform=None`, which means that the constraint is not enforced. See the [GPpyTorch constraints module](#) for additional information.

```
[6]: from botorch.models import SingleTaskGP
from gpytorch.constraints import GreaterThan
from gpytorch.likelihoods import FixedNoiseGaussianLikelihood

# Create a likelihood appropriate for noise-free observations.
# We will use a Gaussian likelihood with a very small variance
    ↪ (noisefree_variance).
# We use a small strictly positive variance instead of zero variance
# to avoid numerical instabilities. If, instead, we set noisefree_variance to
    ↪ 0,
# then gpytorch would issue warning messages and round up to a small positive
    ↪ number.
noisefree_variance = 0.0001
noises = torch.ones(len(train_y)) * noisefree_variance
noise_free_likelihood = FixedNoiseGaussianLikelihood(noise=noises)

# Create our GP model using the training data. By default, BoTorch uses a
    ↪ constant mean and Matern kernel.
model = SingleTaskGP(train_x, train_y, likelihood = noise_free_likelihood)

# Another slightly faster way to do the same thing is:

# from botorch.models import FixedNoiseGP
# model = FixedNoiseGP(torch_train_x, torch_train_y, noises)

# If you want to allow BoTorch to learn the variance of homoscedastic Gaussian
    ↪ noise, drop the likelihood argument as in
# model = SingleTaskGP(train_x, train_y, likelihood = noise_free_likelihood)

model
```

```
[6]: SingleTaskGP(
    (likelihood): FixedNoiseGaussianLikelihood(
      (noise_covar): FixedGaussianNoise()
    )
    (mean_module): ConstantMean()
    (covar_module): ScaleKernel(
      (base_kernel): MaternKernel(
        (lengthscale_prior): GammaPrior()
        (raw_lengthscale_constraint): Positive()
      )
      (outputscale_prior): GammaPrior()
      (raw_outputscale_constraint): Positive()
    )
  )
)
```

You can see from the output above that the GP model has several components: - a likelihood, which describes the likelihood of our observation in terms of the value of the underlying function we want to estimate. We assume that our observations are normally distributed with a mean equal to the function's true value, and a very small variance. One can also assume that we observe the function without noise, but this creates some numerical instabilities in gpytorch. - a mean and covar module, which specifies the mean and kernel function respectively

The covar module has a constraint on its lengthscale saying that it must be positive, and also has a prior that is used when estimating it from data. For now, we will simply use a fixed length scale and so this constraint and prior won't matter for the moment.

Now we change the parameters used in the prior to match what we used in Lab 1 as a default. The constant for the mean because it is 0 by default. Then we can fit the model and see that BoTorch is doing the same thing that we did earlier.

```
[7]: # Because of the way this module is written, we have to construct a neural_
      ↪network parameter
      # from the desired value (0), rather than simply setting the constant to 0.
      model.mean_module.constant = torch.nn.Parameter(torch.tensor(0.))

      # gpytorch only supports these values for the nu parameter for the Matern_
      ↪kernel: 1/2, 3/2, or 5/2
      # If you want to use the value 2, this is supported through the RBF Kernel.
      # (The RBF kernel is the Matern kernel with nu=2).
      model.covar_module.base_kernel.nu = 1.5

      # base_kernel.lengthscale is what we call length_scale in Lab 1 and what we_
      ↪call alpha_1 in the slides.
      # It determines the length scale of the output.
      model.covar_module.base_kernel.lengthscale = 1.

      # model.covar_module.outputscale is what we call alpha0 in our Lab 1 code.
      model.covar_module.outputscale = 1.
```

```
# Train the GP model
model.eval()
```

```
[7]: SingleTaskGP(
    (likelihood): FixedNoiseGaussianLikelihood(
      (noise_covar): FixedGaussianNoise()
    )
    (mean_module): ConstantMean()
    (covar_module): ScaleKernel(
      (base_kernel): MaternKernel(
        (lengthscale_prior): GammaPrior()
        (raw_lengthscale_constraint): Positive()
      )
      (outputscale_prior): GammaPrior()
      (raw_outputscale_constraint): Positive()
    )
  )
```

```
[8]: def plot_posterior(model):

    # Initialize plot
    f, ax = plt.subplots(1, 1, figsize=(6, 4))
    # test model on 101 regular spaced points on the interval [0, 1]
    test_x = torch.linspace(0, 1, 101, dtype=device, device=device)

    with torch.no_grad(): # no need for gradients
        # compute posterior
        posterior = model.posterior(test_x)
        # Get upper and lower confidence bounds (2 standard deviations from the
        mean)
        lower, upper = posterior.mvn.confidence_region()
        # Plot training points as black stars
        ax.plot(train_x.cpu().numpy(), train_y.cpu().numpy(), 'k*')
        # Plot posterior means as blue line
        ax.plot(test_x.cpu().numpy(), posterior.mean.cpu().numpy(), 'b')
        # Shade between the lower and upper confidence bounds
        ax.fill_between(test_x.cpu().numpy(), lower.cpu().numpy(), upper.cpu().
        numpy(), alpha=0.5)

        ax.legend(['Observed Data', 'Mean', 'Credible Interval'])
        plt.tight_layout()

    plot_posterior(model)
```

```
/usr/local/anaconda3/lib/python3.7/site-
packages/gpytorch/lazy/lazy_tensor.py:1810: UserWarning: torch.triangular_solve
```

is deprecated in favor of `torch.linalg.solve_triangular` and will be removed in a future PyTorch release.


`torch.linalg.solve_triangular` has its arguments reversed and does not return a copy of one of the inputs.

```
X = torch.triangular_solve(B, A).solution
```

should be replaced with

```
X = torch.linalg.solve_triangular(A, B). (Triggered internally at  
/Users/runner/work/_temp/anaconda/conda-bld/pytorch_1656352443756/work/aten/src/  
ATen/native/BatchLinearAlgebra.cpp:2189.)
```

```
Linv = torch.triangular_solve(Eye, L, upper=False).solution
```



You should see that this plot matches very closely with the plot that we got in Lab 1 using the default parameters:

```
plot_prediction()
```

Exercise 1 Using BoTorch, reproduce the figure that you generated in Lab 1 with this code

```
plot_prediction(    mean = lambda x : -0.5,    kernel = lambda x1,x2 :  
matern(x1,x2, alpha0=20, length_scales=1.0, nu = 1.5))
```

Exercise 1 Solution

```
[9]: model = SingleTaskGP(train_x, train_y, likelihood = noise_free_likelihood)  
  
model.mean_module.constant = torch.nn.Parameter(torch.tensor(-0.5))
```


```

model.covar_module.base_kernel.nu = 1.5
model.covar_module.base_kernel.lengthscale = 1.
model.covar_module.outputscale = 20.

# Train the GP model
model.eval()

plot_posterior(model)

```



End Exercise 1


2 Choosing the GP hyperparameters

Here we show how to optimize the GP hyperparameters (the lengthscale, output scale, and constant mean of the prior — ν will remain fixed) based on the log marginal likelihood. Here we will show you how to do this manually, explaining all of the steps, and then below we'll show a function that does all of this for you.

2.0.1 Start with a fresh Gaussian process, that doesn't have the parameters we set above

```
[10]: model = SingleTaskGP(train_x, train_y, likelihood = noise_free_likelihood)
model.covar_module.base_kernel.nu = 1.5
model.eval();

plot_posterior(model)
```



Define marginal log likelihood (mll)

```
[11]: from gpytorch.mlls import ExactMarginalLogLikelihood

mll = ExactMarginalLogLikelihood(likelihood=model.likelihood, model=model)

# set mll and all submodules to the numeric data type and device (GPU vs. CPU)
# specified at the start of
# our BoTorch code.
mll = mll.to(train_x)
```

Define optimizer and specify hyperparameters to optimize We will use stochastic gradient descent (`torch.optim.SGD`) to optimize the hyperparameters. In this example, we will use a simple fixed learning rate of 0.2, but in practice the learning rate may need to be adjusted. We will optimize over all parameters in the model, though one could pass a subset of the parameters to the model

to tune only those while leaving the others fixed.

```
[12]: from torch.optim import SGD

optimizer = SGD([{'params': model.parameters()}], lr=0.2)
```

The following code allows us to look at the names of the parameters in the GP regression model that will be optimized. Note that the `outputscale` and `lengthscale` are outputted with a “raw” prepended to them. gpytorch wants to ensure that certain parameters (like the `lengthscale` and `outputscale`) satisfy constraints (specifically, that they are positive). To do this, it maintains a raw version of these parameters that can be any real number, and a transformed version that satisfies the constraint. For example, to satisfy a non-negativity constraint, one can have the transformed version be `exp()` applied to the raw parameter. The parameter `nu` is not included in the parameters that BoTorch optimizes.

```
[13]: print('These are the parameters we will optimize over:')
      for param_name, param in model.named_parameters():
          print(param_name)
```

```
These are the parameters we will optimize over:
mean_module.constant
covar_module.raw_outputscale
covar_module.base_kernel.raw_lengthscale
```

Optimize hyperparameters Now we are ready to write our optimization loop. We will perform 1500 epochs of stochastic gradient descent using our entire training set.

```
[14]: NUM_EPOCHS = 1500

model.train()

for epoch in range(NUM_EPOCHS):
    # clear gradients
    optimizer.zero_grad()
    # forward pass through the model to obtain the output MultivariateNormal
    output = model(train_x)
    # Compute the negative marginal log likelihood
    loss = - mll(output, model.train_targets)
    # Use backward propagation to update the model's
    loss.backward()
    # print every 10 iterations
    if (epoch + 1) % 10 == 0:
        print(
            f"Epoch {epoch+1:>3}/{NUM_EPOCHS} - Negative Marginal Log_
↳ Likelihood : {loss.item():>4.3f} "
            f"lengthscale: {model.covar_module.base_kernel.lengthscale.item():
↳>4.3f} "
            f"constant: {model.mean_module.constant.item():>4.3f} "
```

```
        f"outputscale: {model.covar_module.outputscale.item():>4.3f}"
    )
    optimizer.step()
```

```
Epoch 10/1500 - Negative Marginal Log Likelihood : -0.088 lengthscale: 0.515
constant: -0.184 outputscale: 0.836
Epoch 20/1500 - Negative Marginal Log Likelihood : -0.100 lengthscale: 0.529
constant: -0.305 outputscale: 0.893
Epoch 30/1500 - Negative Marginal Log Likelihood : -0.108 lengthscale: 0.543
constant: -0.390 outputscale: 0.944
Epoch 40/1500 - Negative Marginal Log Likelihood : -0.113 lengthscale: 0.554
constant: -0.452 outputscale: 0.990
Epoch 50/1500 - Negative Marginal Log Likelihood : -0.117 lengthscale: 0.564
constant: -0.500 outputscale: 1.033
Epoch 60/1500 - Negative Marginal Log Likelihood : -0.119 lengthscale: 0.573
constant: -0.537 outputscale: 1.072
Epoch 70/1500 - Negative Marginal Log Likelihood : -0.122 lengthscale: 0.581
constant: -0.567 outputscale: 1.109
Epoch 80/1500 - Negative Marginal Log Likelihood : -0.123 lengthscale: 0.588
constant: -0.591 outputscale: 1.143
Epoch 90/1500 - Negative Marginal Log Likelihood : -0.125 lengthscale: 0.594
constant: -0.611 outputscale: 1.174
Epoch 100/1500 - Negative Marginal Log Likelihood : -0.126 lengthscale: 0.600
constant: -0.628 outputscale: 1.203
Epoch 110/1500 - Negative Marginal Log Likelihood : -0.127 lengthscale: 0.605
constant: -0.643 outputscale: 1.230
Epoch 120/1500 - Negative Marginal Log Likelihood : -0.127 lengthscale: 0.609
constant: -0.656 outputscale: 1.255
Epoch 130/1500 - Negative Marginal Log Likelihood : -0.128 lengthscale: 0.613
constant: -0.667 outputscale: 1.278
Epoch 140/1500 - Negative Marginal Log Likelihood : -0.129 lengthscale: 0.617
constant: -0.677 outputscale: 1.300
Epoch 150/1500 - Negative Marginal Log Likelihood : -0.129 lengthscale: 0.621
constant: -0.686 outputscale: 1.320
Epoch 160/1500 - Negative Marginal Log Likelihood : -0.129 lengthscale: 0.624
constant: -0.694 outputscale: 1.339
Epoch 170/1500 - Negative Marginal Log Likelihood : -0.130 lengthscale: 0.627
constant: -0.701 outputscale: 1.357
Epoch 180/1500 - Negative Marginal Log Likelihood : -0.130 lengthscale: 0.630
constant: -0.707 outputscale: 1.373
Epoch 190/1500 - Negative Marginal Log Likelihood : -0.130 lengthscale: 0.632
constant: -0.713 outputscale: 1.388
Epoch 200/1500 - Negative Marginal Log Likelihood : -0.130 lengthscale: 0.635
constant: -0.719 outputscale: 1.402
Epoch 210/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.637
constant: -0.724 outputscale: 1.416
Epoch 220/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.639
```

constant: -0.728 outputscale: 1.428
 Epoch 230/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.641
 constant: -0.733 outputscale: 1.440
 Epoch 240/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.643
 constant: -0.736 outputscale: 1.450
 Epoch 250/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.644
 constant: -0.740 outputscale: 1.461
 Epoch 260/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.646
 constant: -0.743 outputscale: 1.470
 Epoch 270/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.647
 constant: -0.747 outputscale: 1.479
 Epoch 280/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.648
 constant: -0.749 outputscale: 1.488
 Epoch 290/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.650
 constant: -0.752 outputscale: 1.495
 Epoch 300/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.651
 constant: -0.755 outputscale: 1.503
 Epoch 310/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.652
 constant: -0.757 outputscale: 1.510
 Epoch 320/1500 - Negative Marginal Log Likelihood : -0.131 lengthscale: 0.653
 constant: -0.759 outputscale: 1.516
 Epoch 330/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.654
 constant: -0.761 outputscale: 1.522
 Epoch 340/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.655
 constant: -0.763 outputscale: 1.528
 Epoch 350/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.656
 constant: -0.765 outputscale: 1.534
 Epoch 360/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.656
 constant: -0.767 outputscale: 1.539
 Epoch 370/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.657
 constant: -0.768 outputscale: 1.543
 Epoch 380/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.658
 constant: -0.770 outputscale: 1.548
 Epoch 390/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.658
 constant: -0.771 outputscale: 1.552
 Epoch 400/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.659
 constant: -0.772 outputscale: 1.556
 Epoch 410/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.660
 constant: -0.774 outputscale: 1.560
 Epoch 420/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.660
 constant: -0.775 outputscale: 1.563
 Epoch 430/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.661
 constant: -0.776 outputscale: 1.567
 Epoch 440/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.661
 constant: -0.777 outputscale: 1.570
 Epoch 450/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.662
 constant: -0.778 outputscale: 1.573
 Epoch 460/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.662

constant: -0.779 outputscale: 1.576
 Epoch 470/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.662
 constant: -0.779 outputscale: 1.578
 Epoch 480/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.663
 constant: -0.780 outputscale: 1.581
 Epoch 490/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.663
 constant: -0.781 outputscale: 1.583
 Epoch 500/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.663
 constant: -0.782 outputscale: 1.585
 Epoch 510/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.664
 constant: -0.782 outputscale: 1.587
 Epoch 520/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.664
 constant: -0.783 outputscale: 1.589
 Epoch 530/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.664
 constant: -0.784 outputscale: 1.591
 Epoch 540/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.664
 constant: -0.784 outputscale: 1.593
 Epoch 550/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.665
 constant: -0.785 outputscale: 1.594
 Epoch 560/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.665
 constant: -0.785 outputscale: 1.596
 Epoch 570/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.665
 constant: -0.786 outputscale: 1.597
 Epoch 580/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.665
 constant: -0.786 outputscale: 1.599
 Epoch 590/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.786 outputscale: 1.600
 Epoch 600/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.787 outputscale: 1.601
 Epoch 610/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.787 outputscale: 1.602
 Epoch 620/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.788 outputscale: 1.603
 Epoch 630/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.788 outputscale: 1.604
 Epoch 640/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.788 outputscale: 1.605
 Epoch 650/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.666
 constant: -0.788 outputscale: 1.606
 Epoch 660/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.667
 constant: -0.789 outputscale: 1.607
 Epoch 670/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.667
 constant: -0.789 outputscale: 1.608
 Epoch 680/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.667
 constant: -0.789 outputscale: 1.609
 Epoch 690/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.667
 constant: -0.789 outputscale: 1.609
 Epoch 700/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.667

[illegible]


```

constant: -0.793 outputscale: 1.621
Epoch 1430/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.621
Epoch 1440/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.621
Epoch 1450/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.621
Epoch 1460/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.622
Epoch 1470/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.622
Epoch 1480/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.622
Epoch 1490/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.622
Epoch 1500/1500 - Negative Marginal Log Likelihood : -0.132 lengthscale: 0.669
constant: -0.793 outputscale: 1.622

```

```

[15]: # set model (and likelihood)
      model.eval()

```

```

[15]: SingleTaskGP(
      (likelihood): FixedNoiseGaussianLikelihood(
        (noise_covar): FixedGaussianNoise()
      )
      (mean_module): ConstantMean()
      (covar_module): ScaleKernel(
        (base_kernel): MaternKernel(
          (lengthscale_prior): GammaPrior()
          (raw_lengthscale_constraint): Positive()
          (distance_module): Distance()
        )
        (outputscale_prior): GammaPrior()
        (raw_outputscale_constraint): Positive()
      )
    )

```

```


[16]: plot_posterior(model)
      print(
        f"lengthscale: {model.covar_module.base_kernel.lengthscale.item():>4.3f} "
        f"constant: {model.mean_module.constant.item():>4.3f} "
        f"outputscale: {model.covar_module.outputscale.item():>4.3f} "
        f"nu: {model.covar_module.base_kernel.nu:>4.1f}"
      )

```

```

lengthscale: 0.669 constant: -0.793 outputscale: 1.622 nu: 1.5

```



2.0.2 Here is a faster way to do the same thing

This uses a BoTorch helper function, `fit_gpytorch_model`, that isn't in gpytorch

```
[17]: from botorch.fit import fit_gpytorch_model
model = SingleTaskGP(train_x, train_y, likelihood =
    ↪FixedNoiseGaussianLikelihood(noise=noises))
model.covar_module.base_kernel.nu = 1.5
mll = ExactMarginalLogLikelihood(model.likelihood, model)
fit_gpytorch_model(mll) # This chooses the hyperparameters to maximize the log
    ↪marginal likelihood
```

```
[17]: ExactMarginalLogLikelihood(
    (likelihood): FixedNoiseGaussianLikelihood(
      (noise_covar): FixedGaussianNoise()
    )
    (model): SingleTaskGP(
      (likelihood): FixedNoiseGaussianLikelihood(
        (noise_covar): FixedGaussianNoise()
      )
      (mean_module): ConstantMean()
      (covar_module): ScaleKernel(
        (base_kernel): MaternKernel(
          (lengthscale_prior): GammaPrior()
        )
      )
    )
  )
```

```

        (raw_lengthscale_constraint): Positive()
        (distance_module): Distance()
    )
    (outputscale_prior): GammaPrior()
    (raw_outputscale_constraint): Positive()
)
)
)

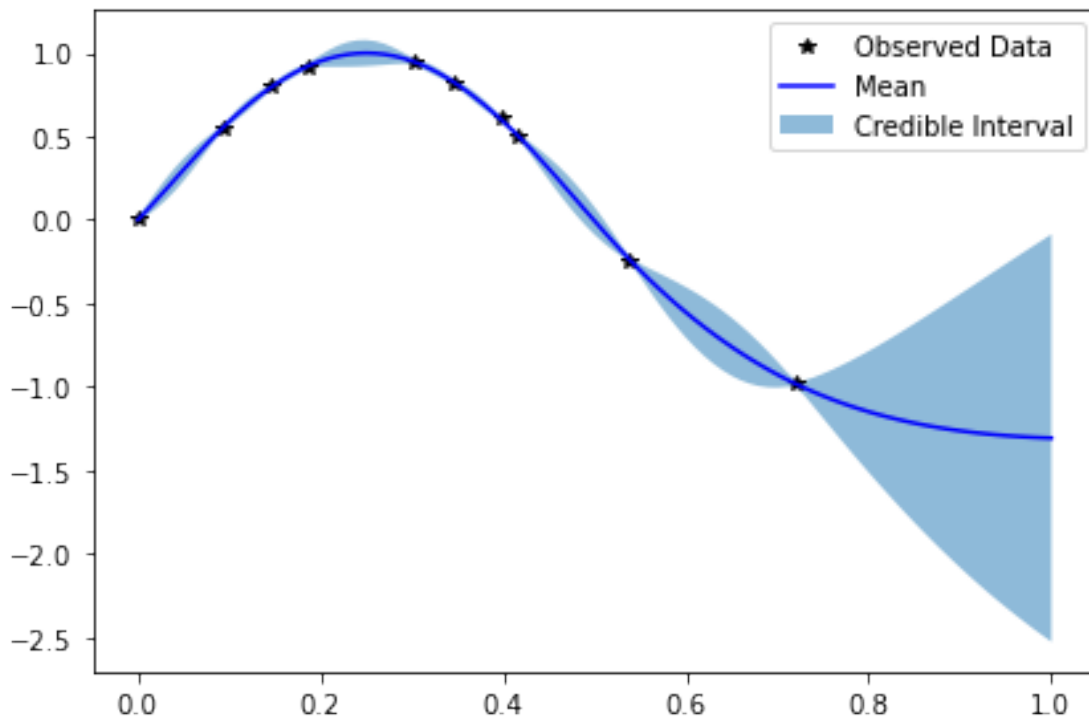
```

```

[18]: plot_posterior(model)
print(
    f"lengthscale: {model.covar_module.base_kernel.lengthscale.item():>4.3f} "
    f"constant: {model.mean_module.constant.item():>4.3f} "
    f"outputscale: {model.covar_module.outputscale.item():>4.3f} "
    f"nu: {model.covar_module.base_kernel.nu:>4.1f}"
)

```

lengthscale: 0.669 constant: -0.793 outputscale: 1.622 nu: 1.5



3 Cross-validation

The following cross-validation plots are for checking model fit.

```
[19]: # Helper function for cross-validation
def remove(T,i):
    # Remove the ith component from the 1-dimensional tensor T
    assert(i<len(T))
    return torch.cat([T[:i], T[i+1:]])

# Check that we handle the two corner cases where i is 0 or len-1
assert(len(remove(train_x,0))==9)
assert(len(remove(train_x,len(train_x)-1))==9)
assert(len(remove(train_x,5))==9)
```

```
[20]: remove(train_y,1)
```

```
[20]: tensor([[ 7.1864e-04],
             [ 7.9687e-01],
             [ 9.2087e-01],
             [ 9.4643e-01],
             [ 8.2510e-01],
             [ 6.0409e-01],
             [ 4.9807e-01],
             [-2.4148e-01],
             [-9.8267e-01]], dtype=torch.float64)
```

```
[21]: def cross_validation(train_x,train_y,nu=1.5):
    loo_mean = []
    loo_sdev = []

    for i in range(len(train_x)):

        # Remove the ith datapoint from the training set
        loo_train_x = remove(train_x,i)
        loo_train_y = remove(train_y,i)
        noisefree_variance = 0.0001
        noises = torch.ones(len(loo_train_y)) * noisefree_variance

        model = SingleTaskGP(loo_train_x, loo_train_y, likelihood = FixedNoiseGaussianLikelihood(noise=noises))
        model.covar_module.base_kernel.nu = nu
        mll = ExactMarginalLogLikelihood(model.likelihood, model)
        fit_gpytorch_model(mll)

        posterior = model.posterior(train_x[i].unsqueeze(0))
        # the posterior mean and variance are 2-d 1x1 tensors.
        # We just need to pull out the single number inside each, which is at
        # index [0][0]
        m = posterior.mean.cpu().detach().numpy()[0][0]
        v = posterior.variance.cpu().detach().numpy()[0][0]
```

```


    loo_mean.append(m)
    loo_sdev.append(np.sqrt(v))

fig, ax = plt.subplots()
train_y_np = train_y.cpu().numpy()[ :,0]

ax.errorbar(train_y_np,loo_mean,loo_sdev,fmt='o')
ax.
↪plot([min(train_y_np),max(train_y_np)], [min(train_y_np),max(train_y_np)], 'k--')
plt.xlabel('observed y')
plt.ylabel('predicted y')

```

```
[22]: cross_validation(train_x,train_y)
```



We want to see that the error bar crosses the dashed line for most of the datapoints (95% of the datapoints), and when it misses it shouldn't miss by too much. This plot looks pretty good.

Exercise 2 Using the function below (which is -1 times a standard test function called Hartmann6), and the 50 training points generated below (uniformly over the unit cube), fit GP regression using a Matern kernel with $\nu=0.5$ and no noise in the observations, Choose the hyperparameters by maximizing the log marginal likelihood.

Then plot a cross-validation plot. You'll see that it looks ok, but some errorbars are far from touching the line. Based on this plot we might want consider another kernel or transforming the

objective.

```
[23]: from botorch.test_functions import Hartmann
      f = Hartmann(negate=True)


      np.random.seed(1)
      train_x = torch.rand(50, 6, device=device, dtype=dtype)
      train_y = f(train_x).unsqueeze(-1)  # add output dimension
```

Exercise 2 Solution

```
[24]: noisefree_variance = 0.0001
      noises = torch.ones(len(train_y)) * noisefree_variance
      model = SingleTaskGP(train_x, train_y, likelihood = FixedNoiseGaussianLikelihood(noise=noises))
      model.covar_module.base_kernel.nu = .5
      mll = ExactMarginalLogLikelihood(model.likelihood, model)
      fit_gpytorch_model(mll)
```

```
[24]: ExactMarginalLogLikelihood(
      (likelihood): FixedNoiseGaussianLikelihood(
        (noise_covar): FixedGaussianNoise()
      )
      (model): SingleTaskGP(
        (likelihood): FixedNoiseGaussianLikelihood(
          (noise_covar): FixedGaussianNoise()
        )
        (mean_module): ConstantMean()
        (covar_module): ScaleKernel(
          (base_kernel): MaternKernel(
            (lengthscale_prior): GammaPrior()
            (raw_lengthscale_constraint): Positive()
            (distance_module): Distance()
          )
          (outputscale_prior): GammaPrior()
          (raw_outputscale_constraint): Positive()
        )
      )
    )
```

```
[25]: cross_validation(train_x, train_y)
```



[]: